


Parallel Particle Filter

Kelvin Silva

Introduction

Particle Filter:

Monte Carlo Localization



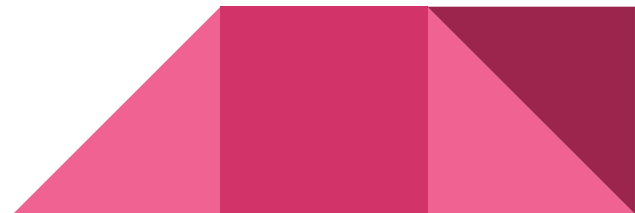
Serial = Slow

1 Motion Update

2 Measurement Update

3 Resampling step

For $N \rightarrow$ big number doing these 3 steps iteratively is slow



Parallel = Faster

Can use GPU to do Motion, Measurement updates on every particle at the same time



Resampling is Harder

Resampling step is most commonly done with iterative approach based on a cumulative sum of weights.

So need to use Metropolis resampling



Parallel:

Pseudocode for Metropolis resampling.

METROPOLIS-ANCESTORS($\mathbf{w} \in [0, \infty)^N$, $B \in \{1, 2, \dots\}$) $\rightarrow \{1, \dots, N\}^N$

```
1  for each  $i \in \{1, \dots, N\}$ 
2       $k \leftarrow i$ 
3      for  $n = 1, \dots, B$ 
4           $u \sim \mathcal{U}[0, 1]$ 
5           $j \sim \mathcal{U}\{1, \dots, N\}$ 
6          if  $u \leq w^j / w^k$ 
7               $k \leftarrow j$ 
8   $a^i \leftarrow k$ 
9  return  $\mathbf{a}$ 
```

Serial:

```
# particle filter resampling
p3 = []
index = int(random.random() * N)
beta = 0.0
mw = max(w)
for i in range(N):
    beta += random.random() * 2.0 * mw
    while beta > w[index]:
        beta -= w[index]
        index = (index + 1) % N
    p3.append(p[index])
p = p3
```

Code Review

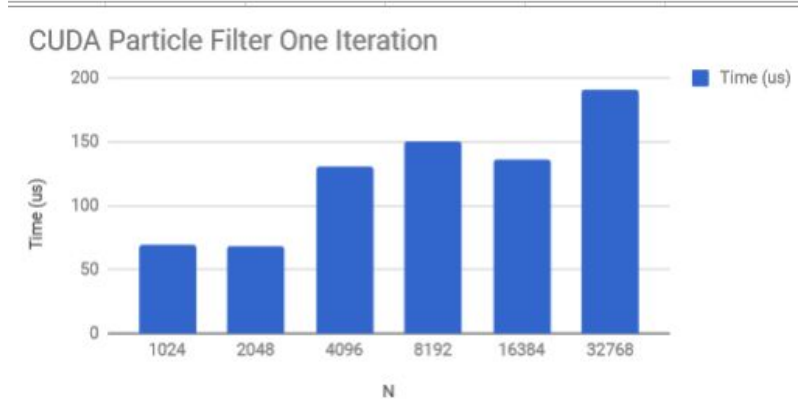
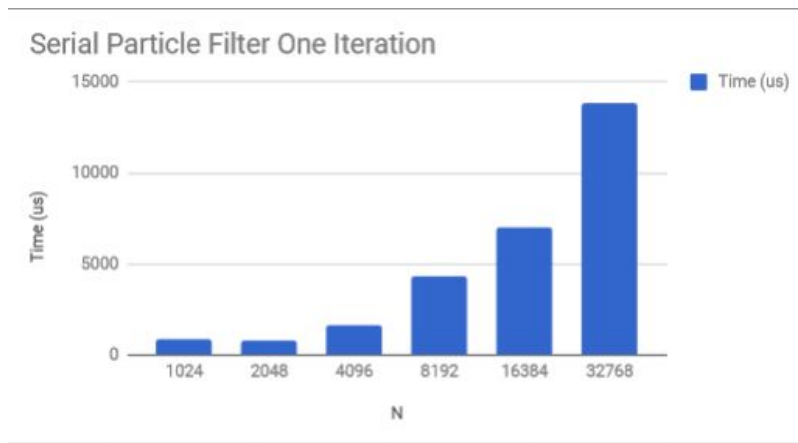
<https://github.com/kelvinsilva/ams148-gpu-programming-cuda/tree/master/particle-filter-parallel>

Thrun Udacity:

<https://gist.github.com/kelvinsilva/6c19c12c54c5d873aad01f65017cb7d9>



Results Speedup



Conclusion

For large amount of particles better to parallelize

Due to algorithm runtime complexity, even an unoptimized parallel particle filter will run faster than a serial one for LARGE N

For small amounts (less than 1024) iterative is better

GPU Memory limiting factor

